



AI研习社

CS231n Assignment3


王煦中

[知乎专栏：喵神大人的深度工坊](#)

- **Part 1 Image Captioning with Vanilla RNNs**
- **Part 2 Image Captioning with LSTMs**
- **Part 3 Network Visualization**
- **Part 4 Style Transfer**
- **Part 5 Generative Adversarial Networks(GAN)**



Part 1 Image Captioning with Vanilla RNNs



➤ Vanilla RNN 单元(前向):

- (右箭头) $h_{\{t\}}$ 从左到右的传播过程 $h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$
- (上箭头) $h_{\{t\}}$ temporal softmax 输出层 $y_t = \text{softmax}(W_y * h_t + b_y)$

```
next_h = np.tanh(x.dot(Wx) + prev_h.dot(Wh) + b)
cache = (x, prev_h, Wh, Wx, b, next_h)
```



Part 1 Image Captioning with Vanilla RNNs

$$h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$$

➤ Vanilla RNN 单元(反向): 上游导数记为dout

- $db = \text{sum}(\mathbf{dout} * (1 - \tanh^2(x)))$
- $dW_x = \mathbf{dout} * (1 - \tanh^2(x)) \text{ dot } x$
- $dW_h = \mathbf{dout} * (1 - \tanh^2(x)) \text{ dot } h$
- $dx = \mathbf{dout} * (1 - \tanh^2(x)) \text{ dot } W_x$
- $dh = \mathbf{dout} * (1 - \tanh^2(x)) \text{ dot } W_h$

```
(x, prev_h, Wh, Wx, b, next_h) = cache
dtheta = dnext_h * (1 - next_h ** 2) #(N,H)
db = np.sum(dtheta, axis=0)
dWx = x.T.dot(dtheta) #(D,N) * (N,H)
dWh = prev_h.T.dot(dtheta) #(H,N) * (N,H)
dprev_h = dtheta.dot(Wh.T) # (N,H) * (H,H)= (N,H)
dx = dtheta.dot(Wx.T) # (N,H)*(H,D)=(N,D)
```



AI研习社

Part 1 Image Captioning with Vanilla RNNs

➤完整 Vanilla RNN (前向):

- 对RNN单元循环T次，T是序列的长度 $h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$


```
prev_h = h0
h = np.zeros((N, T, H))
for t in range(T):
    xt = x[:, t, :]
    next_h, cache = rnn_step_forward(xt, prev_h, Wx, Wh, b)
    h[:, t, :] = next_h
    prev_h = next_h
```



Part 1 Image Captioning with Vanilla RNNs

➤完整 Vanilla RNN (反向):

- 一部分来自右边记为dprev_h
- 一部分来自上面记为dh
- 首先逆序: `for t in range(T-1, -1, -1):`



- 对于每一个时间片t, 上面来的导数是 $dh[:, t, :]$ (形状是(N, T, H)), 对于**最后一个单元**, 它的右边没有传来导数, 所以**初始化dprev_h是0**。
- 于是对于每一个RNN单元, 需要传进去的导数就是 $dh[:, t, :] + dprev_h$ 。

```
cache = (xt, prev_h, Wh, Wx, b, next_h)
```

```
dnext_h = dh[:, t, :] + dprev_h #上面+右边, 初始右边=0
```

```
dx[:, t, :], dprev_h, dWxt, dWht, dbt = rnn_step_backward(dnext_h, cache)
```

- dWx , dWh , db 和 dh 这四个参数在**不同时刻是共享的**, 因此需要在**每一个时刻把它们都加起来做更新**。

```
dWx += dWxt #不同时刻共享的
```

```
dWh += dWht
```

```
db += dbt
```

```
dh0 = dprev_h
```

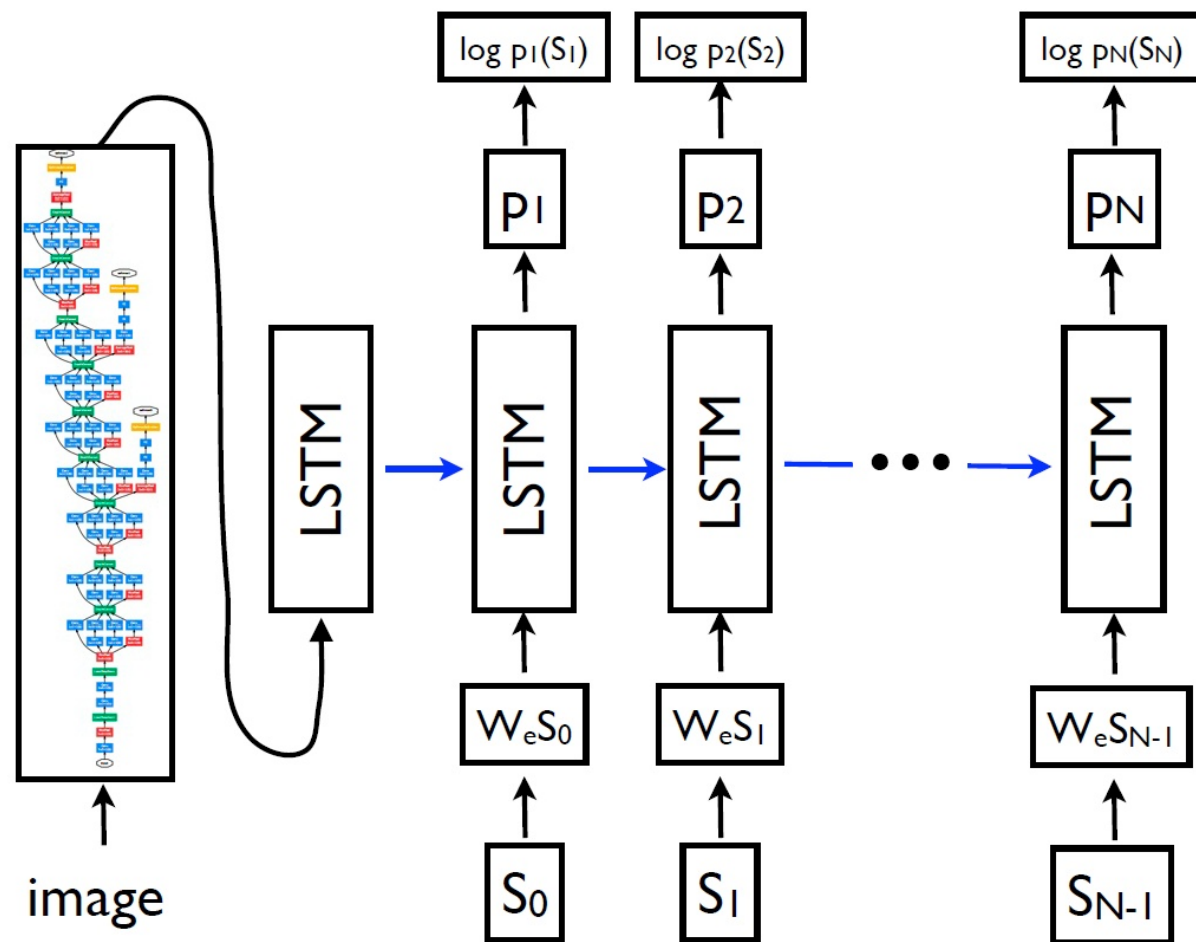


Part 1 Image Captioning with Vanilla RNNs

➤ Image Captioning

- 隐藏层：从训练好的vgg16的fc7层中取出特征，当做 h_0 输入到RNN中
- 单词的表示：
- Word embedding, 把onehot编码的单词映射为一个向量
- Word embedding(反向)：类似relu、dropout的反向

```
dW = np.zeros(W.shape)
np.add.at(dW, x, dout)#按照x的下标给出dout
```



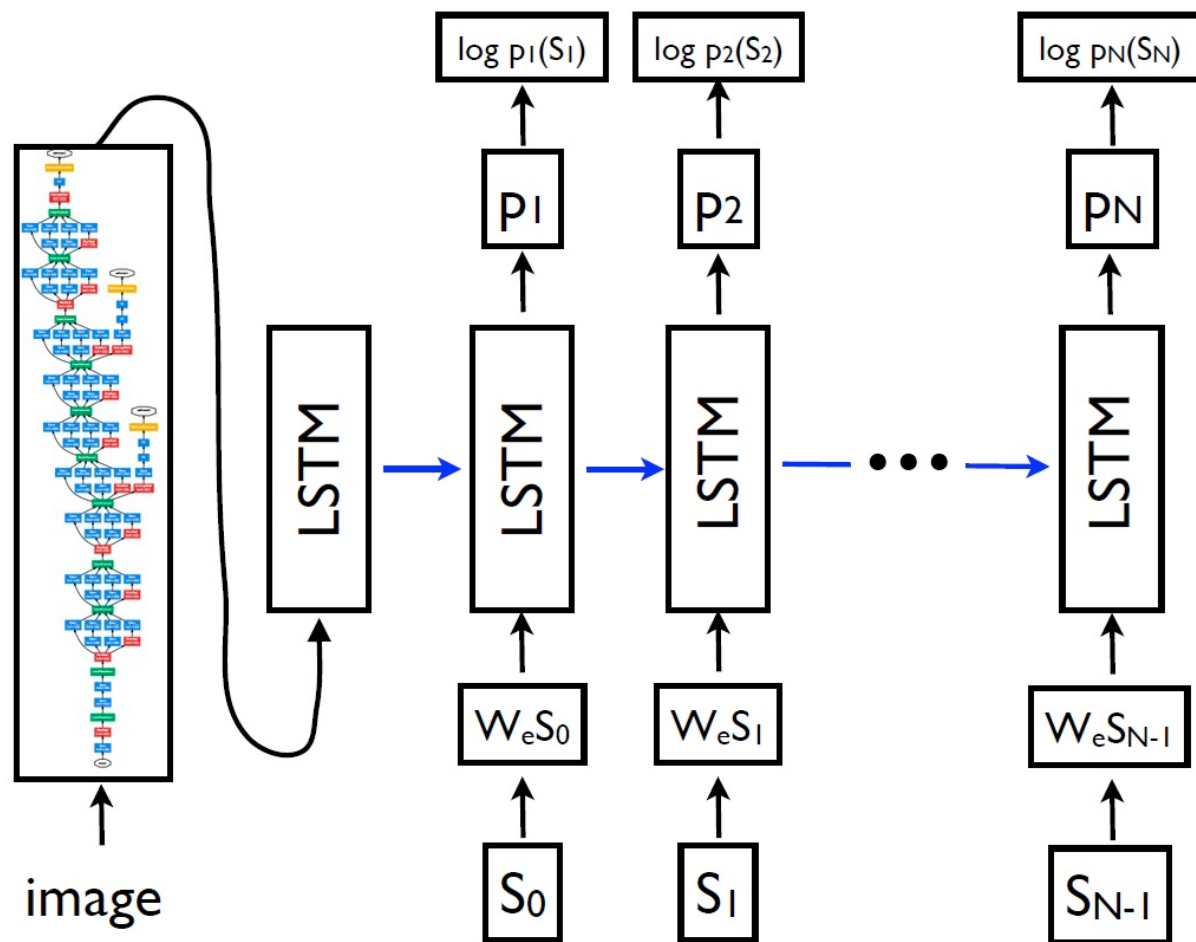


Part 1 Image Captioning with Vanilla RNNs

➤ Image Captioning 训练

- 下一个词当做当前词的label。captions的第一个词是start标记，最后一个词是end标记。

```
# forward
# affine hidden state(N,D) -> (N,H)
affine, cache_affine = affine_forward(features, W_proj, b_proj)
# word embedding
embed, cache_embed = word_embedding_forward(captions_in, W_embed)
# rnn
if self.cell_type == 'rnn':
    h, cache_rnn = rnn_forward(embed, affine, Wx, Wh, b)
elif self.cell_type == 'lstm':
    h, cache_rnn = lstm_forward(embed, affine, Wx, Wh, b)
# temporal_affine
h_affine, cache_h_affine = temporal_affine_forward(h, W_vocab, b_vocab)
# softmax
loss, dx = temporal_softmax_loss(h_affine, captions_out, mask)
```





Part 1 Image Captioning with Vanilla RNNs

➤ Image Captioning 测试

- 构造一个输入，它的第一个单词是start标记。和训练时不同的是，测试时我们不知道T的大小，因此限定了最大输出长度

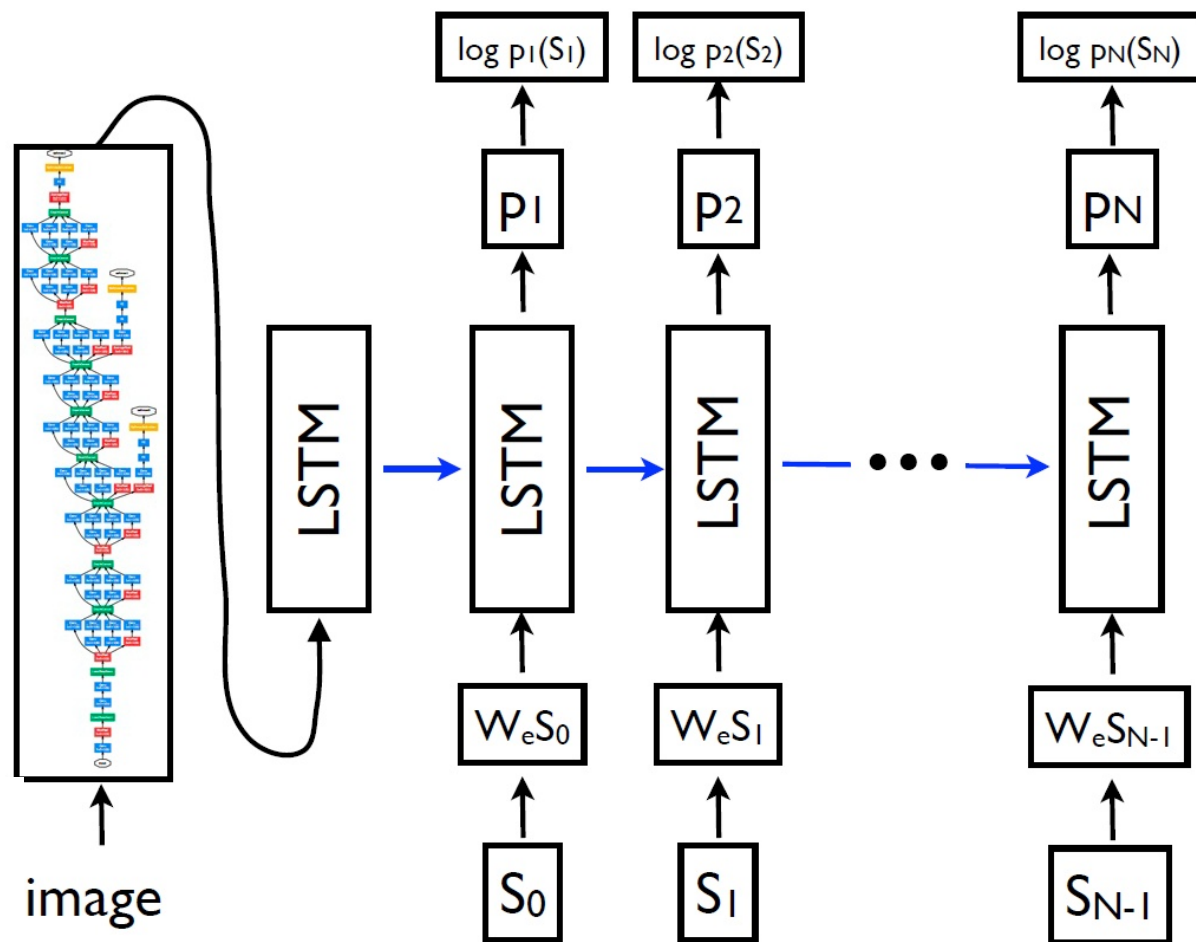
```
# for each step
for t in range(1,max_length):
    # word embed
    embed, cache_embed = word_embedding_forward(x, W_embed)

    # rnn step
    if self.cell_type == 'rnn':
        next_h, cache = rnn_step_forward(embed, prev_h, Wx, Wh, b)
    prev_h = next_h

    # affine
    h_affine, cache_h_affine = affine_forward(next_h, W_vocab, b_vocab)


    # max
    x = np.argmax(h_affine, axis=1)

    captions[:,t] = x
```



Part 2 Image Captioning with LSTMs

AI研习社



➤ LSTM 单元(前向):

- (上右箭头) $c_{\{t\}}$
- (下右箭头) $h_{\{t\}}$
- (上箭头) $h_{\{t\}}$ softmax
- f、i、g、o的线性部分可以通过以一次计算完成

$$f_t = \text{sigmoid}(W_{fx}x_t + W_{fh}h_{t-1} + b_f)$$

$$i_t = \text{sigmoid}(W_{ix}x_t + W_{ih}h_{t-1} + b_i)$$

$$g_t = \tanh(W_{gx}x_t + W_{gh}h_{t-1} + b_g)$$

$$c_t = f_t \circ c_{t-1} + i_t \circ g_t$$

$$o_t = \text{sigmoid}(W_{ox}x_t + W_{oh}h_{t-1} + b_o)$$

$$h_t = o_t \circ \tanh(c_t)$$


$$z = x.\text{dot}(Wx) + \text{prev_h}.\text{dot}(Wh) + b$$

Part 2 Image Captioning with LSTMs

AI研习社

➤ LSTM 单元(反向):

- 从右侧传回来的值有两个，分别是dnext_h和dnext_c
- 四个线性部分记作i, f, o, g



$$f_t = \text{sigmoid}(W_{fx}x_t + W_{fh}h_{t-1} + b_f)$$

$$i_t = \text{sigmoid}(W_{ix}x_t + W_{ih}h_{t-1} + b_i)$$

$$g_t = \tanh(W_{gx}x_t + W_{gh}h_{t-1} + b_g)$$

$$c_t = f_t \circ c_{t-1} + i_t \circ g_t$$

$$o_t = \text{sigmoid}(W_{ox}x_t + W_{oh}h_{t-1} + b_o)$$

$$h_t = o_t \circ \tanh(c_t)$$

$$df = (dnext_h + dh2c) * prev_c * f * (1 - f)$$

$$di = (dnext_h + dh2c) * g * i * (1 - i)$$

$$dg = (dnext_h + dh2c) * i * (1 - g**2)$$

$$dprev_c = (dnext_h + dh2c) * f$$

$$do = dnext_h * \tanh(c) * o * (1 - o)$$

$$dh2c = dnext_h * o * (1 - \tanh(c)**2)$$

```
d = np.hstack((di, df, do, dg)) #(N, 4H)
```


Part 2 Image Captioning with LSTMs

AI研习社

➤完整 LSTM (前向):

- 比RNN多了一个c, 初始为0

```
prev_h = h0
h = np.zeros((N, T, H))
prev_c = np.zeros((N, H))
cache = {}
for t in range(T):
    xt = x[:, t, :]
    next_h, next_c, cache[t] = lstm_step_forward(xt, prev_h, prev_c, Wx, Wh, b)
    h[:, t, :] = next_h
    prev_h = next_h
    prev_c = next_c
```



Part 2 Image Captioning with LSTMs

AI研习社

➤完整 LSTM (反向):

- 比RNN多了一个c

- 首先逆序: `for t in range(T-1, -1, -1):`

- 对于每一个时间片t, 上面来的导数是 $dh[:, t, :]$ (形状是(N, T, H)), 对于最后一个单元, 它的右边没有传来导数, 所以初始化 $dprev_h$ 与 $dnext_c$ 都是0。

- 于是对于每一个LSTM单元, 需要传进去的导数就是 $dh[:, t, :] + dprev_h$ 。

```
dnext_h = dh[:, t, :] + dprev_h #上面+右边, 初始右边=0
```

```
dx[:, t, :], dprev_h, dprev_c, dWxt, dWht, dbt = lstm_step_backward(dnext_h, dnext_c, cache[t])
```

```
dnext_c = dprev_c
```


- dWx , dWh , db 和 dh 这四个参数在不同时刻是共享的, 因此需要在每一个时刻把它们都加起来做更新。

```
dWx += dWxt #不同时刻共享的
```

```
dWh += dWht
```

```
db += dbt
```

```
dh0 = dprev_h
```



Part 3 Network Visualization


AI研习社

➤ Saliency maps

hay




0



quail




1



Tibetan mastiff




2




Border terrier brown bear, bruin, Ursus arctos



3



4





Part 3 Network Visualization

AI研习社

➤ Fooling Images

- 调整图片，使其骗过我们训练好的网络
- 也就是用目标分类对输入图片的梯度来迭代更新输入图片



Part 3 Network Visualization

AI研习社

➤ Class Visualization

- 网络关注的每个类别
分类的梯度


- 首先计算梯度

```
loss = tf.nn.softmax_cross_entropy_loss
loss -= l2_reg * tf.nn.l2_loss(model.parameters)
grad = tf.gradients(loss, model.parameters)
```

- 迭代，更新噪声图片

```
clip = tf.clip_by_norm(grad, clip_norm)
X = sess.run(clip)
```

tarantula
Iteration 100 / 100



上更新目标

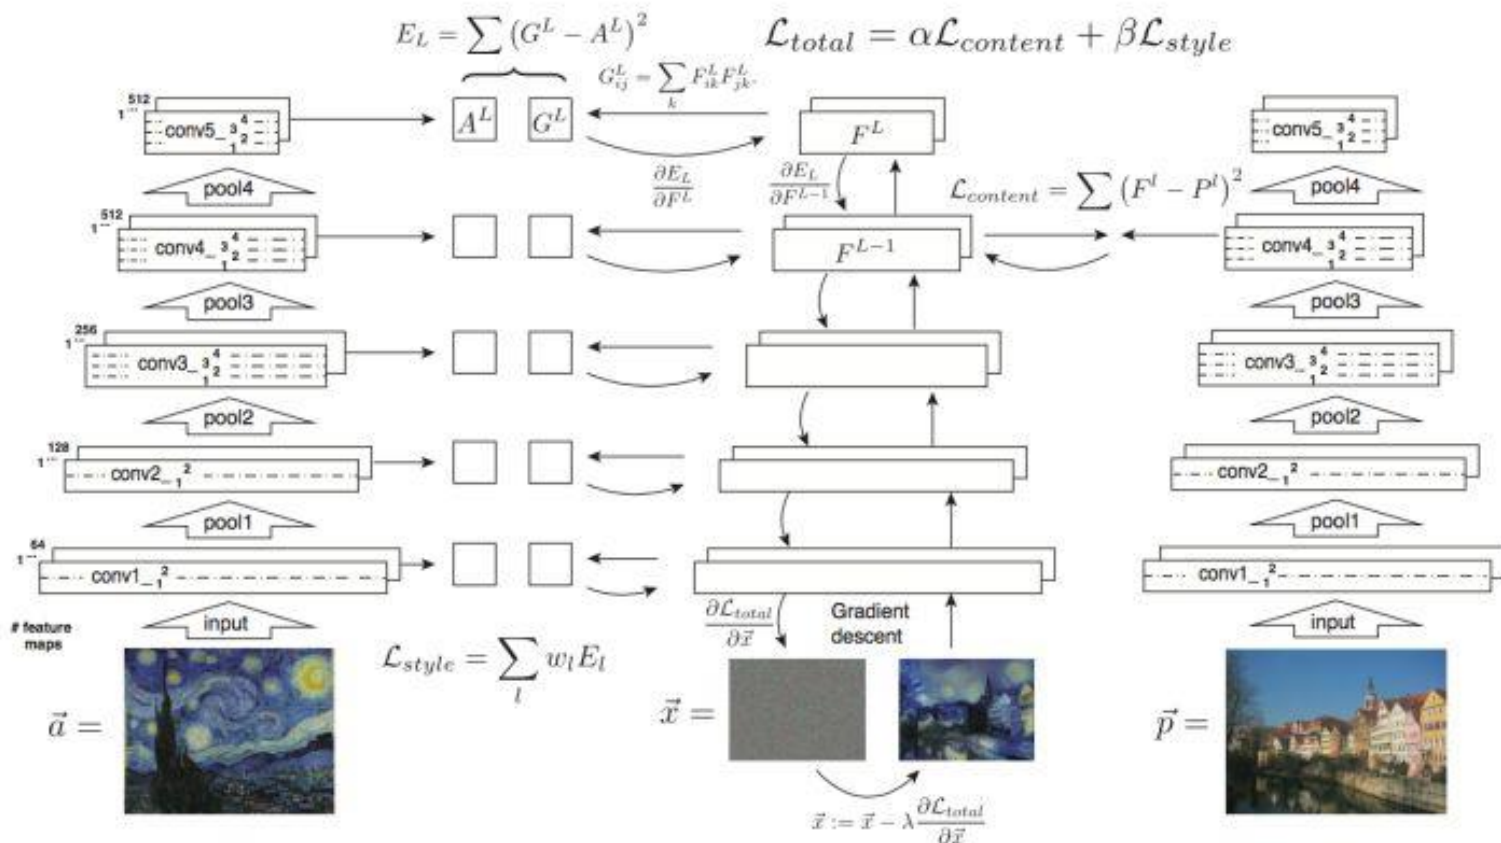
```
s=model.classifier) # scalar
e, same size as model.image
```




Part 4 Style Transfer

➤ 网络结构

- 定义一个**新的loss**，然后对**随机噪声图片**进行梯度更新。
- 注意**三个图片**是经过**同样的训练好的网络的**。





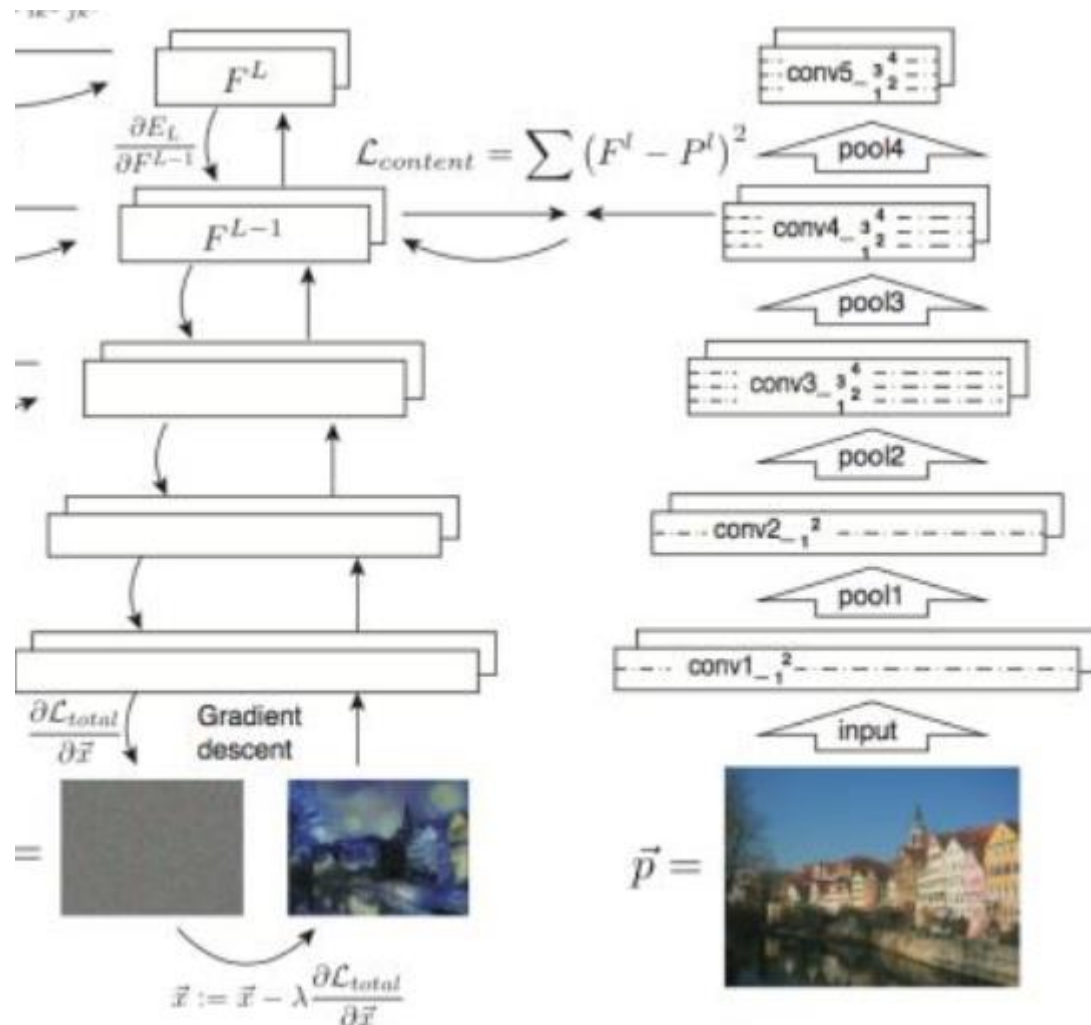
Part 4 Style Transfer

➤ content loss

- 内容loss反映了生成的图片内容和源内容图片的差异

$$L_c = w_c * \sum (F_{i,j}^l - P_{i,j}^l)^2$$

- 其中 $F_{i,j}^l$ 是生成图片在网络中第 l 层的feature map, $P_{i,j}^l$ 是源内容图片在网络中第 l 层的feature map





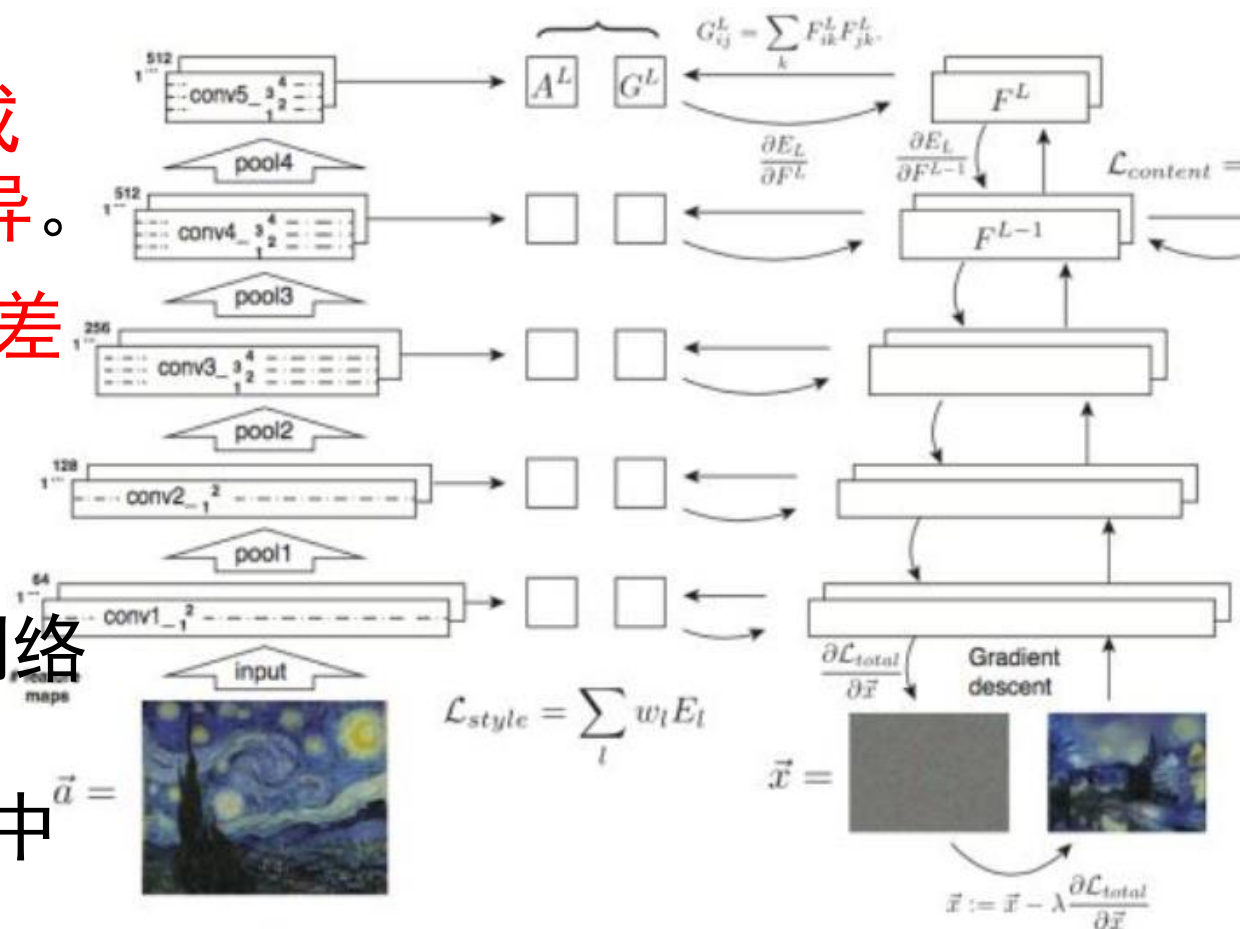
Part 4 Style Transfer

➤ style loss

- 风格loss目的就是为了**衡量生成的图片风格和源风格图片的差异**。
- 而风格是用feature map的**协方差矩阵**度量。 $G = F^T * F$

$$L_s = w_s * \sum (G_{i,j}^l - A_{i,j}^l)^2$$

- 其中 $G_{i,j}^l$ 是**生成图片**在网络中第 l 层的feature map, $A_{i,j}^l$ 是**源风格图片**在网络中第 l 层的feature map





Part 4 Style Transfer

➤ total variation loss

- 使生成图片更平滑

$$L_{tv} = w_t * \sum_{c=1}^3 \sum_{i=1}^{H-1} \sum_{j=1}^{W-1} (x_{i,j+1,c} - x_{i,j,c})^2 + (x_{i+1,j,c} - x_{i,j,c})^2$$

- 这个式子是在H和W维度上计算相邻像素的差值的平方和

```
pixel_dif1 = img[:, 1:, :W-1, :] - img[:, :H-1, :W-1, :]  
pixel_dif2 = img[:, :H-1, 1:, :] - img[:, :H-1, :W-1, :]
```

- tf里提供了绝对值版本的total variation loss：
tf.image.total_variation()。也就是不算平方，只用绝对值。



AI研习社

Part 4 Style Transfer

- 最后3个loss加在一起就行了
- 可以调的参数：loss的**三个权重**，内容loss和风格loss的**层**选择
- 当风格loss权重为0时，是对源内容图像的重建

Content Source Img.



Style Source Img.





AI研习社

Part 5 Generative Adversarial Networks

➤ GAN

- GAN其实就是**生成器**和**判别器**之间的博弈，一方面**判别器**从真实数据和生成数据中不断学习**提高自己的判别能力**，另一方面**生成器**不断迭代以提高自己的**欺骗能力**。

$$\text{minimize}_G \text{ maximize}_D \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(x))]$$

- 1. **更新生成器**使得判别器做出**正确判断**的概率**下降**
- 2. **更新判别器**使得判别器做出**正确判断**的概率**上升**

$$\text{maximize}_D \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(x))]$$

- 1. 等价于 1.* **更新生成器**使得判别器做出**错误判断**的概率**上升**

$$\text{maximize}_G \mathbb{E}_{x \sim p_z(x)} [\log D(G(x))]$$



AI研习社

Part 5 Generative Adversarial Networks

➤ Vanilla Gan

- Discriminator

- Fully connected layer from size 784 to 256
- LeakyReLU with alpha 0.01
- Fully connected layer from 256 to 256
- LeakyReLU with alpha 0.01
- Fully connected layer from 256 to 1

- Generator

- Fully connected layer from `tf.shape(z)[1]` (the number of noise dimensions) to 1024
- ReLU
- Fully connected layer from 1024 to 1024
- ReLU
- Fully connected layer from 1024 to 784
- TanH (To restrict the output to be $[-1, 1]$)



Part 5 Generative Adversarial Networks

➤ Vanilla Gan Loss: sigmoid_cross_entropy loss

$$\text{maximize}_D \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(x))]$$

$$\text{maximize}_G \mathbb{E}_{x \sim p_z(x)} [\log D(G(x))]$$

```
with tf.variable_scope("") as scope:
    #scale images to be -1 to 1
    logits_real = discriminator(preprocess_img(x))
    # Re-use discriminator weights on new inputs
    scope.reuse_variables()
    logits_fake = discriminator(G_sample)
    D_label = tf.ones_like(logits_real)
    G_fake_label = tf.zeros_like(logits_fake)
    G_real_label = tf.ones_like(logits_fake)
    D_loss1 = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(logits = logits_real, labels = D_label))
    D_loss2 = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(logits = logits_fake, labels = G_fake_label))
    G_loss = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(logits = logits_fake, labels = G_real_label))
    return D_loss1 + D_loss2, G_loss
```




Part 5 Generative Adversarial Networks

➤ Least Squares GAN

- Loss由交叉熵改变为Least Squares:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

```
D_loss = tf.reduce_mean(tf.square(score_real - 1)) + tf.reduce_mean(tf.square(score_fake))  
G_loss = tf.reduce_mean(tf.square(score_fake - 1))  
return 0.5 * D_loss, 0.5 * G_loss
```



AI研习社

Part 5 Generative Adversarial Networks

➤ Deep Convolutional GANs

• Discriminator

- 32 Filters, 5x5, Stride 1, Leaky ReLU(alpha=0.01)
- Max Pool 2x2, Stride 2
- 64 Filters, 5x5, Stride 1, Leaky ReLU(alpha=0.01)
- Max Pool 2x2, Stride 2
- Flatten
- Fully Connected size 4 x 4 x 64, Leaky ReLU(alpha=0.01)
- Fully Connected size 1

• Generator

- Fully connected of size 1024, ReLU
- BatchNorm
- Fully connected of size 7 x 7 x 128, ReLU
- BatchNorm
- Resize into Image Tensor
- 64 conv2d^T (transpose) filters of 4x4, stride 2, ReLU
- BatchNorm
- 1 conv2d^T (transpose) filter of 4x4, stride 2, TanH

Part 5

AI研习社



➤ WGAN-G

- Discriminator:

- 64 Filters of 4x4
- 128 Filters of 4x4
- BatchNorm
- Flatten
- Fully connecte
- Fully connecte

- Generator:

- Loss: grad
- <https://www>



works



AI研习社

课后讨论

那么多GAN，究竟哪个更好？

=====><https://zhuanlan.zhihu.com/p/31563676>



AI研习社

CS231n Assignment3 Fin